

CSCI3150 Introduction to Operating Systems

Lecture 11: Memory Management Part II: Paging

Hong Xu

<https://github.com/henryhxu/CSCI3150>

Overview

- ▣ Paging
- ▣ Translation lookaside buffer, TLB
- ▣ Advanced page tables


Paging: Introduction

Concept of Paging


- ▣ Segmentation divides an address space into **variable** sizes of segments
- ▣ Paging splits up an address space into **fixed-sized** units, or **pages**
- ▣ With paging, **physical memory** is also split into some number of pages called a **page frame**

Example: A Simple Paging Scheme

- 128-byte physical memory with 16-byte page frames
- 64-byte address space with 16-byte pages



A Simple 64-byte Address Space



64-Byte Address Space Placed In Physical Memory

Advantages of Paging


- ▣ **Flexibility:** Supporting the abstraction of address space effectively
 - ◆ Don't need assumptions on how heap and stack grow and are used.
- ▣ **Simplicity:** ease of free-space management
 - ◆ The page in address space and the page frame are the same size.
 - ◆ Easy to allocate and keep a free list

Page table

- A **page table** is needed to translate the virtual address to physical address
 - ◆ So we know where each virtual page is in the physical memory!
 - ◆ In our example, page 0 is in physical frame 3, page 1 in frame 7, page 2 in frame 5, and page 3 in frame 2
- Page table is **per process, obviously!**


0	reserved for OS	page frame 0 of physical memory
16	(unused)	page frame 1
32	page 3 of AS	page frame 2
48	page 0 of AS	page frame 3
64	(unused)	page frame 4
80	page 2 of AS	page frame 5
96	(unused)	page frame 6
112	page 1 of AS	page frame 7
128		

Example: Page Table in Kernel Physical Memory




Address Translation

- Two components in the virtual address
 - VPN: virtual page number
 - Offset: offset within the page



- Example: virtual address 21 in a 64-byte address space




A screenshot of the 'User Login' form for OpenVPN Access Server. It includes fields for 'Username' and 'Password', each with a small icon to its left. Below the fields is a dark blue 'Sign In' button.

(not this one...)

Example: Address Translation

- The virtual address 21 in 64-byte address space



Where are Page Tables Stored?

- ▣ Page tables can get awfully large!
 - ◆ 32-bit address space with 4-KB pages, 20 bits for VPN
 - Page offset for 4 Kbyte pages: 12 bits
 - 2^{20} entries, $\sim 1M$, 4 Bytes per page table entry $\Rightarrow$ 4MB per process!
- ▣ Page tables for each process are stored in memory


What is in the Page Table?

- ▣ The page table is a **data structure** that is used to map the virtual address to physical address.
 - ◆ Simplest form: a linear page table, an array
- ▣ The OS **indexes** the array by VPN, and looks up the page-table entry (PTE)

Common Flags of Page Table Entry

- ▣ **Valid Bit:** Indicating whether the particular translation is valid (unused space)
- ▣ **Protection Bit:** Indicating whether the page could be read from, written to, or executed from
- ▣ **Present Bit:** Indicating whether this page is in physical memory or on disk (swapped out)
- ▣ **Dirty Bit:** Indicating whether the page has been modified since it was brought into memory
- ▣ **Reference Bit (Accessed Bit):** Indicating that a page has been accessed

Example: x86 Page Table Entry



An x86 Page Table Entry(PTE)

- ❑ P: present
- ❑ R/W: read/write bit
- ❑ U/S: supervisor
- ❑ A: accessed bit
- ❑ D: dirty bit
- ❑ PFN: the page frame number

Paging: Also Too Slow

- ▣ To find a location of the desired PTE, the **starting location** of the page table is **needed**.
- ▣ For every memory reference, paging requires the OS to perform one **extra** memory reference

Accessing Memory With Paging


```
1      // Extract the VPN from the virtual address
2      VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4      // Form the address of the page-table entry (PTE)
5      PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7      // Fetch the PTE
8      PTE = AccessMemory(PTEAddr)
9
```


Accessing Memory With Paging

```
10 // Check if process can access the page
11 if (PTE.Valid == False)
12     RaiseException(SEGMENTATION_FAULT)
13 else if (CanAccess(PTE.ProtectBits) == False)
14     RaiseException(PROTECTION_FAULT)
15 else
16     // Access is OK: form physical address and fetch it
17     offset = VirtualAddress & OFFSET_MASK
18     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19     Register = AccessMemory(PhysAddr)
```

Translation Lookaside Buffer (TLB)

- ▣ Part of the chip's memory-management unit (MMU).
- ▣ A hardware cache of **popular** virtual-to-physical address translation.



TLB Basic Algorithms

```
1: VPN = (VirtualAddress & VPN_MASK ) >> SHIFT
2: (Success, TlbEntry) = TLB_Lookup(VPN)
3: if (Success == TRUE) { // TLB Hit
4:     if (CanAccess (TlbEntry.ProtectBit) == True ) {
5:         offset = VirtualAddress & OFFSET_MASK
6:         PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:         AccessMemory (PhysAddr)
8:     }else RaiseException (PROTECTION_ERROR)
```

TLB Basic Algorithms (Cont.)

```
11:      }else{ //TLB Miss
12:          PTEAddr = PTBR + (VPN * sizeof(PTE))
13:          PTE = AccessMemory(PTEAddr)
14:          if(PTE.Valid == False)
15:              RaiseException(SEGFAULT) ;
16:          else{
17:              TLB_Insert( VPN , PTE.PFN , PTE.ProtectBits)
18:              RetryInstruction()
19:          }
```

- ◆ (11-12 lines) The hardware accesses the page table to find the translation.
- ◆ (16 lines) updates the TLB with the translation.

Example: Accessing An Array

- How a TLB can improve performance

VPN	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:  int sum = 0 ;
1:  for(i=0; i<10; i++) {
2:      sum += a[i] ;
3:  }
```

First access to a page is a miss, but the subsequent ones are hits


3 misses and 7 hits.
Thus TLB hit rate is 70%.

**The TLB improves performance
due to spatial locality**

Locality


□ Temporal Locality

- ◆ An instruction or data item that has been recently accessed will likely be accessed soon again



□ Spatial Locality

- ◆ If a program accesses memory at address x , it will likely soon access memory near x .



Who Handles a TLB Miss?

- ▣ Hardware handles the TLB misses entirely on CISC
 - ◆ The hardware has to know exactly where the page tables are located in memory.
 - ◆ The hardware would “walk” the page table, find the correct page-table entry and extract the desired translation, update and retry instruction.
 - ◆ **hardware-managed TLB.**
 - ◆ **Intel x86**
- ▣ RISC has what is known as a **software-managed TLB**
 - ◆ On a TLB miss, the hardware raises an exception (trap handler)
 - **Trap handler is code** within the OS that is written with the express purpose of handling TLB miss.

TLB Control Flow algorithm (OS Handled)

```
1:      VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2:      (Success, TlbEntry) = TLB_Lookup(VPN)
3:      if (Success == True) // TLB Hit
4:      if (CanAccess(TlbEntry.ProtectBits) == True)
5:          Offset = VirtualAddress & OFFSET_MASK
6:          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:          Register = AccessMemory(PhysAddr)
8:      else
9:          RaiseException(PROTECTION_FAULT)
10:     else // TLB Miss
11:         RaiseException(TLB_MISS)
```


TLB entry

- ▣ TLB is managed by **fully-associative*** method.
 - ◆ A typical TLB has 32, 64, or 128 entries.
 - ◆ Hardware searches the entire TLB in parallel to find the desired translation.
 - ◆ Other bits: valid bits, protection bits, **address-space identifier**, dirty bit




*More on cache associativity: <http://csillustrated.berkeley.edu/PDFs/handouts/cache-3-associativity-handout.pdf>


TLB Issue: Context Switching



TLB Issue: Context Switching



TLB Issue: Context Switching




TLB Table

VPN	PFN	valid	prot
10	100	1	rwX
-	-	-	-
10	170	1	rwX
-	-	-	-

Can't Distinguish which entry is meant for which process

Solution

- Add an address space identifier (ASID) field in the TLB



Another Case

▣ Two processes share a page

- ◆ Process 1 is sharing physical frame 101 with Process 2.
- ◆ P1 maps this frame into the 10th page of its address space.
- ◆ P2 maps this frame to the 50th page of its address space.

VPN	PFN	valid	prot	ASID
10	101	1	rwX	1
-	-	-	-	-
50	101	1	rwX	2
-	-	-	-	-

Sharing of pages is **useful** as it reduces the number of physical pages/frames in use.


On the side: ASID vs. PID

- ❑ ASID seems very similar to PID right?
 - ◆ To distinguish processes
 - ◆ They are related, yes
- ❑ Can you just use PID then?
 - ◆ Obviously no (otherwise it won't be called ASID...)
 - ◆ ASID (hardware) uses much fewer bits than PID (OS)
- ❑ How to use ASID when the PID space is much bigger?
 - ◆ <https://stackoverflow.com/questions/52813239/how-many-bits-there-are-in-a-tlb-asid-tag-for-intel-processors-and-how-to-handl>
 - ◆ Caveat: ASID is only needed when the process is active (from TLB's perspective)

TLB Replacement Policy

▣ Least Recently Used, LRU


- ◆ Evict an entry that has not recently been used
- ◆ Take advantage of *locality* in the memory-reference stream.



Total 11 TLB miss

A Real TLB Entry

All 64 bits of this TLB entry (example of MIPS R4000)




Flag	Content
19-bit VPN	The rest reserved for the kernel.
24-bit PFN	Systems can support with up to 64GB of main memory(pages).
Global bit(G)	Used for pages that are globally-shared among processes.
ASID	OS can use to distinguish between address spaces.
Coherence bit(C)	determine how a page is cached by the hardware.
Dirty bit(D)	marking when the page has been written.
Valid bit(V)	tells the hardware if there is a valid translation present in the entry.

Advanced Page Tables

Paging: Linear Tables

- We usually have one page table for every process in the system.
 - ◆ Assume 32-bit address space with 4KB pages and 4-byte page-table entry.




$$\text{Page table size} = \frac{2^{32}}{2^{12}} * 4\text{Byte} = 4\text{MByte}$$

Page tables are **too big** and thus consume too much memory.

Paging: Smaller Tables

- Page tables are too big and thus consume too much memory.
 - ◆ Assume that 32-bit address space with **16KB** pages and 4-byte page-table entry.




$$\frac{2^{32}}{2^{14}} * 4 = 1MB \text{ per page table}$$

Big pages lead to internal fragmentation.

Problem

- Single page table for the address space of process.




A 16KB Address Space with 1KB Pages

PFN	valid	prot	present	dirty
10	1	r-x	1	0
-	0	-	-	-
-	0	-	-	-
-	0	-	-	-
15	1	rw-	1	1
...	...	...	...	...
-	0	-	-	-
3	1	rw-	1	1
23	1	rw-	1	1

A Page Table For 16KB Address Space

Problem

- Most of the page table is **unused**, full of invalid entries.




PFN	valid	prot	present	dirty
9	1	r-x	1	0
-	0	-	-	-
-	0	-	-	-
-	0	-	-	-
15	1	rw-	1	1
...	...	...	...	...
-	0	-	-	-
3	1	rw-	1	1
23	1	rw-	1	1

A Page Table For 16KB Address Space

Hybrid Approach: Paging and Segments

- Page table for each segment
 - ◆ The base register for each of these segments contains the physical address of a linear page table for that segment.
 - ◆ The bound register: indicate the end of the page table.
- Example: Each process has **three** page tables associated with it.



32-bit virtual address space with 4KB pages

Seg value	Content
00	unused segment
01	code
10	heap
11	stack

TLB Miss on Hybrid Approach

- ▣ The hardware gets **physical address** from **page table**.
 - ◆ The hardware uses the segment bits (SN) to determine which base-and-bounds pair to use.
 - ◆ The hardware then takes the **physical address** therein and **combines** it with the VPN as follows to form the address of the page table entry (PTE)

```
01:    SN = (VirtualAddress & SEG_MASK) >> SN_SHIFT
```

```
02:    VPN = (VirtualAddress & VPN_MASK) >> VPN_SHIFT
```

```
03:    AddressOfPTE = Base[SN] + (VPN * sizeof(PTE))
```



Problem of the Hybrid Approach

- ▣ Hybrid approach is not without problems.
 - ◆ If we have a large but sparsely-used heap, we can still end up with a lot of page table waste.
 - ◆ Causing external fragmentation, again!

Multi-Level Page Tables

- ▣ Turn the linear page table into something like a tree.
 - ◆ Chop up the page table into page-sized units.
 - ◆ If an entire page of page-table entries is invalid, don't allocate that page of the page table at all.
 - ◆ To track whether a page of the page table is valid, use a new structure, called **page directory**.

Multi-level Page Tables



Multi-level Page Tables

▣ Page Directory

- ◆ The page directory contains one entry per page of the page table.
- ◆ It consists of a number of **entries, i.e. page directory entries (PDE)**.
- ◆ A PDE has a valid bit and page frame number (PFN).

▣ Advantage

- ◆ Only allocates page-table space in proportion to the amount of address space you are using.
- ◆ The OS can grab the next free page when it needs to allocate or grow a page table.

▣ Disadvantage

- ◆ Multi-level table is a small example of a **time-space trade-off**.
- ◆ **Complexity**.

On-the-side: Indirection

- ▣ Does this remind you of something familiar?
- ▣ Adding another table for a table => adding another level of indirection

All problems in computer science can be solved by another level of indirection.

- David Wheeler

Example


0000 0000	code
0000 0001	code
...	(free)
	(free)
0000 0100	heap
0000 0101	heap
.....	(free)
	(free)
1111 1110	stack
1111 1111	stack

- ▣ Page 0,1: code
- ▣ Page 4,5: heap
- ▣ Page 254, 255: stack

Example

Flag	Detail
Address space	16 KB (2^{14} Byte)
Page size	64 B
Virtual address	14 bit
VPN	8 bit
Offset	6 bit
Page table entry	4 Byte


A 16-KB Address Space With 64-byte Pages



▣ Single level paging

- ◆ 256 page table entries: 2^8 entries
- ◆ Page table size: $256 * 4 \text{ Byte} = 1 \text{ Kbyte}$
- ◆ Page table needs 16 pages of physical memory (64B each): $1024/64 = 16$

Example: single level page table




Example: two level paging

▣ Page directory index


- ◆ The page table has 16 pages.
- ◆ 16 entries for page directory: one entry per page of the page table.
- ◆ $16 \times 4B = 64B$ is required for page directory. → One page
- ◆ 4 bits for page directory index.

$$\text{PDEAddr} = \text{PageDirBase} + (\text{PDIndex} * \text{sizeof(PDE)})$$



- ▣ If the page-directory entry is **invalid**, raise an exception (The access is invalid).

Example: two level paging




Example: Page Table Index

▣ Page table index

- ◆ It is used to find the address of the page table entry.

$$\text{PTEAddr} = (\text{PDE.PFN} \ll \text{SHIFT}) + (\text{PTIndex} * \text{sizeof}(\text{PTE}))$$




Examples

Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
101	1	—	0	—	45	1	rw-

Single level paging: 16 pages

→ Two level paging: 3 pages



Multi-level Page Table Control Flow

```
01:     VPN = (VirtualAddress & VPN_MASK) >> SHIFT
02:     (Success, TlbEntry) = TLB_Lookup(VPN)
03:     if (Success == True)           //TLB Hit
04:         if (CanAccess(TlbEntry.ProtectBits) == True)
05:             Offset = VirtualAddress & OFFSET_MASK
06:             PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
07:             Register = AccessMemory(PhysAddr)
08:         else RaiseException(PROTECTION_FAULT);
09:     else // perform the full multi-level lookup
```

- ◆ (1 line) extract the virtual page number (VPN)
- ◆ (2 lines) check if the TLB holds the translation for this VPN
- ◆ (5-8 lines) extract the page frame number from the relevant TLB entry, and form the desired physical address and access memory

Multi-level Page Table Control Flow

```
11:         else
12:             PDIndex = (VPN & PD_MASK) >> PD_SHIFT
13:             PDEAddr = PDBR + (PDIndex * sizeof(PDE))
14:             PDE = AccessMemory(PDEAddr)
15:             if(PDE.Valid == False)
16:                 RaiseException(SEGMENTATION_FAULT)
17:             else // PDE is Valid: now fetch PTE from PT
```

- ◆ (11 lines) extract the Page Directory Index (PDIndex)
- ◆ (13 lines) get Page Directory Entry (PDE)
- ◆ (15-17 lines) Check PDE valid flag. If valid flag is true, fetch Page Table entry from Page Table

The Translation Process: Remember the TLB

```
18:     PTIndex = (VPN & PT_MASK) >> PT_SHIFT
19:     PTEAddr = (PDE.PFN << SHIFT) + (PTIndex * sizeof(PTE))
20:     PTE = AccessMemory(PTEAddr)
21:     if(PTE.Valid == False)
22:         RaiseException(SEGMENTATION_FAULT)
23:     else if(CanAccess(PTE.ProtectBits) == False)
24:         RaiseException(PROTECTION_FAULT);
25:     else
26:         TLB_Insert(VPN, PTE.PFN , PTE.ProtectBits)
27:         RetryInstruction()
```


Inverted Page Tables

- ▣ Keeping a single page table that has an entry for each physical page of the system.
- ▣ The entry tells us which process is using this page, and which virtual page of that process maps to this physical page.