

CSCI3150 Introduction to Operating Systems

Lecture 15: File Systems Part III: Log-structured File System

Hong Xu

<https://github.com/henryhxu/CSCI3150>

Outline

- ▣ Key Idea: Writing Sequentially
- ▣ Indirect Mapping and Checkpoint Region
- ▣ Directories
- ▣ Garbage Collection
- ▣ Crash Recovery

Overview

- In the early 90's, a new file system, the log-structured file system (LFS) was developed
- Motivation
 - ◆ Memory sizes were growing.
 - ◆ Large gap between random IO and sequential IO performance.
 - ◆ Existing file systems perform poorly on common workloads.
- In this chapter, we study Log-Structured Filesystem (LFS).
 - ◆ How can a file system **transform all writes into sequential writes**?

John Ousterhout




Mendel Rosenblum


Key Idea: Writing to Disk Sequentially

- How do we transform all updates to file-system state into a series of sequential writes to disk?

- data update




- metadata needs to be updated too. (Ex. inode)



Writing Sequentially, and Effectively!

- ▣ Writing to the disk sequentially is not enough to guarantee efficient writes
 - ◆ Disk may rotate between the writes
 - ◆ Write buffering
 - ◆ Segment: a set of sequential writes that are written to the disk with a single unit.
 - ◆ Keep track of updates in **memory buffer**. (a few MB)
 - ◆ Write them to disk all at once



Issue #1: How Much to Buffer

- Time to write D MB

$$T_{write} = T_{position} + \frac{D}{R_{peak}}$$

- Effective write bandwidth

$$R_{effective} = \frac{D}{T_{write}} = \frac{D}{T_{position} + \frac{D}{R_{peak}}}$$

- We'd like to make the effective write bandwidth close to peak bandwidth with some fraction F ($0 < F < 1$)

$$R_{effective} = \frac{D}{T_{position} + \frac{D}{R_{peak}}} = F \times R_{peak}$$

Issue #1: How Much to Buffer

- Then, D can be computed as follows

$$D = F \times R_{peak} \times (T_{position} + \frac{D}{R_{peak}})$$

$$D = (F \times R_{peak} \times T_{position}) + \left(F \times R_{peak} \times \frac{D}{R_{peak}} \right)$$

$$D = \frac{F}{1-F} \times R_{peak} \times T_{position}$$

- Example: Positioning time 10ms, peak transfer rate 100MByte/sec, we like to achieve 90% of the peak rate

$$D = 0.9 \times 0.1 \times 100 \text{ Mbyte/sec} \times 0.01 \text{ secs} = 9 \text{ Mbyte}$$


- What is D if F = 0.95?

Outline

- ▣ Key Idea: Writing Sequentially
- ▣ Indirect Mapping and Checkpoint Region
- ▣ Directories
- ▣ Garbage Collection
- ▣ Crash Recovery


Issue #2: How to Find Inodes?

- ▣ The position of the inodes keep **changing**.
- ▣ The Inode Map
 - ◆ A data structure that contains the location of the most recent inode for a given inode number.
 - ◆ Places the chunk of updated inode map right next to the updated inode (**why?**)
 - ◆ Where to find the inode map?



Issue #2: How to Find Inodes?

- How to find the inode map spread across the disk?
 - The LFS must have a fixed location on disk to begin a file lookup.
- Checkpoint Region**
 - A fixed location in the LFS partition.
 - Contains the pointers to the latest of the inode map.
- Wait, then every update still needs to seek to this CR!
 - Inode map also helps to solve the recursive update problem (coming up soon)



Reading a file from the disk

- ▣ Reading a file block
 - ◆ Read a checkpoint region
 - ◆ Read inode map
 - ◆ Read inode
 - ◆ Read data block
- ▣ What about sequential read?
 - ◆ It may become random read.


LFS is optimized for writes.

Outline

- ▣ Key Idea: Writing Sequentially
- ▣ Indirect Mapping and Checkpoint Region
- ▣ Directories
- ▣ Garbage Collection
- ▣ Crash Recovery

Issue #3: What About Directories?

- ▣ Directory: a set of <inode, filename>
- ▣ How does LFS store directory data?
- ▣ Creating a file: `foo`
 - ◆ Update the directory inode. (inode #: `dir`)
 - ◆ Update the directory entry. (`foo`, `k`)
 - ◆ Update inode for the created file. (inode #: `k`)
 - ◆ Update the data block for the created file.



Issue #3: What About Directories?

- ❑ Recursive update: A serious problem in any FS that never updates in place
 - ◆ Whenever an inode is updated, its **location** on disk changes
 - To keep track of inodes, a directory may have to record a collection of (name, inode-**location**) pairs instead
 - ◆ This would have also entailed **recursive updates** to the directory that points to the file, the parent of that directory, ..., all the way up the file system tree
- ❑ LFS cleverly avoids this problem with imap
 - ◆ The directory is still (name, inode-**num**) pairs.
 - ◆ The imap keeps the inode-**num** to inode-**location** mappings
 - The change in inode location is never reflected in the directory itself


Outline

- ▣ Key Idea: Writing Sequentially
- ▣ Indirect Mapping and Checkpoint Region
- ▣ Directories
- ▣ Garbage Collection
- ▣ Crash Recovery


Issue #4: Garbage Collection

- LFS keeps writing newer versions of a file.
- Garbage: LFS leaves the older versions of file structures all over the disk.
- An example of garbage

- ◆ Overwrite the data block:



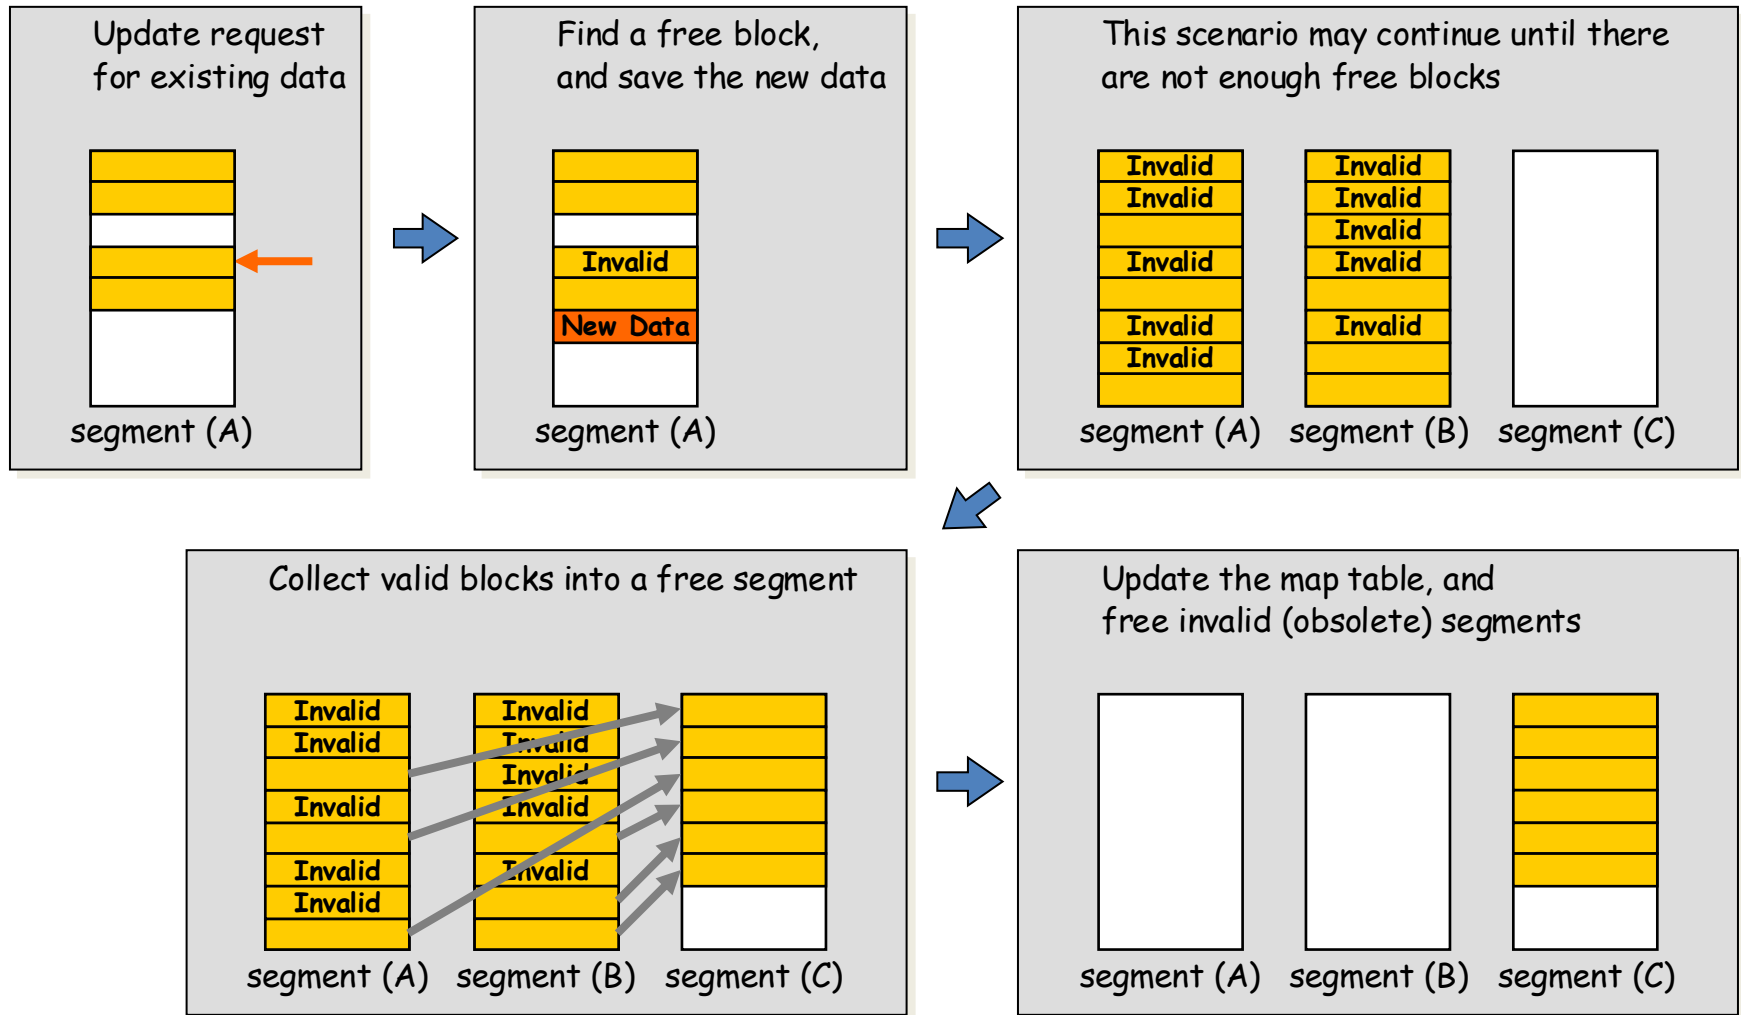
- ◆ Append a block to that original file k:



Issue #4: Garbage Collection


- ▣ What to do with the older versions of the block
 - ◆ Versioning filesystem: keep the old blocks and allow the users to restore to the older version of the filesystem status.
 - ◆ LFS: periodically clean the older versions of the file data, inodes and other structures.
- ▣ Unit of garbage collection: Segment
 - ◆ Read a number of old segments, M segments.
 - ◆ Identify the valid blocks.
 - ◆ Write them to a number of new segments (in memory), N segments.
 - ◆ Write N segments to the disk.
 - ◆ Then, $N < M$.

Issue #4: Garbage Collection



Mechanism: Segment Summary Block

- Store the **inode number** and the **offset for each data block** in it.
- In garbage collection, we need to identify the obsolete blocks.
- Compare the block address of file k, block offset 0, based upon the Segment Summary and based upon the in-memory imap. If they coincide, the block is alive



Policy of Garbage Collection

- ▣ When to clean
 - ◆ Periodically
 - ◆ When a system is idle
 - ◆ When the disk is full
- ▣ Which block to consolidate? (heuristic from the original LFS paper)
 - ◆ Hot segment: the blocks are updated periodically
 - ◆ Cold segment: the blocks are not updated.
 - ◆ Hot segment: clean later.
 - ◆ Cold segment: clean sooner.
- ▣ Other policies are possible

Outline


- ▣ Key Idea: Writing Sequentially
- ▣ Indirect Mapping and Checkpoint Region
- ▣ Directories
- ▣ Garbage Collection
- ▣ Crash Recovery

Crash recovery

- What if the crash happens when the LFS is in the middle of writing the segment to the disk?
- LFS maintains a set of segments as a linked list in memory




- LFS organizes the filesystem partition as a log (a linked list of the segments).
 - Two checkpoint regions: one at the beginning and the one at the end.



Crash recovery

□ Consistent Update on CR

- ◆ Write timestamp at the beginning of CR.
- ◆ Write CR body.
- ◆ Write time stamp at the end of the CR.
- ◆ When crash occurs, chooses the most recent CR with valid consistent time stamps.



□ Crash recovery

- ◆ Read the CR and rebuild imap.
- ◆ Perform roll-forward.
 - Start from the first segment in CR.
 - Scan the valid segment following the "next segment" pointer and update the imap.

Summary

- ▣ Introduce a new approach to updating the disk.
 - ◆ **Shadow paging** in database system, **Copy-on-Write** in file system.
- ▣ Gather all updates into an in-memory segment.
 - ◆ Write them out together sequentially.
- ▣ LFS-style is excellent for performance on many different devices.
 - ◆ Hard drives, parity-based RAIDs, even Flash-based SSDs.
- ▣ Some modern commercial filesystems adopt a similar copy-on-write approach even though it generates garbage.
 - ◆ NetApp's **WAFL**, Sun's **ZFS** and Linux **btrfs**
 - ◆ In particular, WAFL turns cleaning problem into a feature, by providing old versions of the file system via **snapshots**.